

[1]: import pandas as pd

[2]: df = pd.read_csv("HousingData.csv")

[3]: df.head()

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LSTAT | MEDV |
|---|---------|------|-------|------|-------|-------|------|--------|-----|-----|---------|--------|-------|------|
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1 | 296 | 15.3 | 396.90 | 4.98 | 24.0 |
| 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242 | 17.8 | 396.90 | 9.14 | 21.6 |
| 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2 | 242 | 17.8 | 392.83 | 4.03 | 34.7 |
| 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222 | 18.7 | 394.63 | 2.94 | 33.4 |
| 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222 | 18.7 | 396.90 | NaN | 36.2 |

[4]: df.shape

(506, 14)

[5]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    CRIM        486 non-null    float64
1    ZN          486 non-null    float64
2    INDUS       486 non-null    float64
3    CHAS        486 non-null    float64
4    NOX         506 non-null    float64
5    RM          506 non-null    float64
6    AGE         486 non-null    float64
7    DIS         506 non-null    float64
8    RAD         506 non-null    int64
9    TAX         506 non-null    int64
10   PTRATIO     506 non-null    float64
11   B           506 non-null    float64
12   LSTAT       486 non-null    float64
13   MEDV        506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

```
[6]: df.describe()
```

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LSTAT |
|-------|------------|------------|------------|------------|------------|------------|------------|------------|------------|------------|------------|------------|------------|
| count | 486.000000 | 486.000000 | 486.000000 | 486.000000 | 506.000000 | 506.000000 | 486.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 |
| mean | 3.611874 | 11.211934 | 11.083992 | 0.069959 | 0.554695 | 6.284634 | 68.518519 | 3.795043 | 9.549407 | 408.237154 | 18.455534 | 356.674032 | 12.715432 |
| std | 8.720192 | 23.388876 | 6.835896 | 0.255340 | 0.115878 | 0.702617 | 27.999513 | 2.105710 | 8.707259 | 168.537116 | 2.164946 | 91.294864 | 7.155871 |
| min | 0.006320 | 0.000000 | 0.460000 | 0.000000 | 0.385000 | 3.561000 | 2.900000 | 1.129600 | 1.000000 | 187.000000 | 12.600000 | 0.320000 | 1.730000 |
| 25% | 0.081900 | 0.000000 | 5.190000 | 0.000000 | 0.449000 | 5.885500 | 45.175000 | 2.100175 | 4.000000 | 279.000000 | 17.400000 | 375.377500 | 7.125000 |
| 50% | 0.253715 | 0.000000 | 9.690000 | 0.000000 | 0.538000 | 6.208500 | 76.800000 | 3.207450 | 5.000000 | 330.000000 | 19.050000 | 391.440000 | 11.430000 |
| 75% | 3.560263 | 12.500000 | 18.100000 | 0.000000 | 0.624000 | 6.623500 | 93.975000 | 5.188425 | 24.000000 | 666.000000 | 20.200000 | 396.225000 | 16.955000 |
| max | 88.976200 | 100.000000 | 27.740000 | 1.000000 | 0.871000 | 8.780000 | 100.000000 | 12.126500 | 24.000000 | 711.000000 | 22.000000 | 396.900000 | 37.970000 |

```
[7]: # Data Cleaning
df=df.fillna(df.mean())
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    CRIM        506 non-null    float64
1    ZN          506 non-null    float64
2    INDUS       506 non-null    float64
3    CHAS        506 non-null    float64
4    NOX         506 non-null    float64
5    RM          506 non-null    float64
6    AGE         506 non-null    float64
7    DIS         506 non-null    float64
8    RAD         506 non-null    int64
9    TAX         506 non-null    int64
10   PTRATIO     506 non-null    float64
11   B           506 non-null    float64
12   LSTAT       506 non-null    float64
13   MEDV        506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

```
[8]: corr=df.corr() #Find the correlation (relationship) between each column in the DataFrame
corr
```

[8]:

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LSTAT | MEDV |
|---------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| CRIM | 1.000000 | -0.182930 | 0.391161 | -0.052223 | 0.410377 | -0.215434 | 0.344934 | -0.366523 | 0.608886 | 0.566528 | 0.273384 | -0.370163 | 0.434044 | -0.379695 |
| ZN | -0.182930 | 1.000000 | -0.513336 | -0.036147 | -0.502287 | 0.316550 | -0.541274 | 0.638388 | -0.306316 | -0.308334 | -0.403085 | 0.167431 | -0.407549 | 0.365943 |
| INDUS | 0.391161 | -0.513336 | 1.000000 | 0.058035 | 0.740965 | -0.381457 | 0.614592 | -0.699639 | 0.593176 | 0.716062 | 0.384806 | -0.354597 | 0.567354 | -0.478657 |
| CHAS | -0.052223 | -0.036147 | 0.058035 | 1.000000 | 0.073286 | 0.102284 | 0.075206 | -0.091680 | 0.001425 | -0.031483 | -0.109310 | 0.050055 | -0.046166 | 0.179882 |
| NOX | 0.410377 | -0.502287 | 0.740965 | 0.073286 | 1.000000 | -0.302188 | 0.711461 | -0.769230 | 0.611441 | 0.668023 | 0.188933 | -0.380051 | 0.572379 | -0.427321 |
| RM | -0.215434 | 0.316550 | -0.381457 | 0.102284 | -0.302188 | 1.000000 | -0.241351 | 0.205246 | -0.209847 | -0.292048 | -0.355501 | 0.128069 | -0.602962 | 0.695360 |
| AGE | 0.344934 | -0.541274 | 0.614592 | 0.075206 | 0.711461 | -0.241351 | 1.000000 | -0.724353 | 0.449989 | 0.500589 | 0.262723 | -0.265282 | 0.574893 | -0.380223 |
| DIS | -0.366523 | 0.638388 | -0.699639 | -0.091680 | -0.769230 | 0.205246 | -0.724353 | 1.000000 | -0.494588 | -0.534432 | -0.232471 | 0.291512 | -0.483429 | 0.249929 |
| RAD | 0.608886 | -0.306316 | 0.593176 | 0.001425 | 0.611441 | -0.209847 | 0.449989 | -0.494588 | 1.000000 | 0.910228 | 0.464741 | -0.444413 | 0.468440 | -0.381626 |
| TAX | 0.566528 | -0.308334 | 0.716062 | -0.031483 | 0.668023 | -0.292048 | 0.500589 | -0.534432 | 0.910228 | 1.000000 | 0.460853 | -0.441808 | 0.524545 | -0.468536 |
| PTRATIO | 0.273384 | -0.403085 | 0.384806 | -0.109310 | 0.188933 | -0.355501 | 0.262723 | -0.232471 | 0.464741 | 0.460853 | 1.000000 | -0.177383 | 0.373343 | -0.507787 |
| B | -0.370163 | 0.167431 | -0.354597 | 0.050055 | -0.380051 | 0.128069 | -0.265282 | 0.291512 | -0.444413 | -0.441808 | -0.177383 | 1.000000 | -0.368886 | 0.333461 |
| LSTAT | 0.434044 | -0.407549 | 0.567354 | -0.046166 | 0.572379 | -0.602962 | 0.574893 | -0.483429 | 0.468440 | 0.524545 | 0.373343 | -0.368886 | 1.000000 | -0.721975 |
| MEDV | -0.379695 | 0.365943 | -0.478657 | 0.179882 | -0.427321 | 0.695360 | -0.380223 | 0.249929 | -0.381626 | -0.468536 | -0.507787 | 0.333461 | -0.721975 | 1.000000 |

[9]:

df.rename(columns={'MEDV': 'PRICE'}, inplace=True)

[10]:

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
Column Non-Null Count Dtype
--- -
0 CRIM 506 non-null float64
1 ZN 506 non-null float64
2 INDUS 506 non-null float64
3 CHAS 506 non-null float64
4 NOX 506 non-null float64
5 RM 506 non-null float64
6 AGE 506 non-null float64
7 DIS 506 non-null float64
8 RAD 506 non-null int64
9 TAX 506 non-null int64
10 PTRATIO 506 non-null float64
11 B 506 non-null float64
12 LSTAT 506 non-null float64
13 PRICE 506 non-null float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB

```
[11]: x=df.drop('PRICE',axis=1) # Independent Columns
      y=df['PRICE'] #Target Column Price
```

```
[12]: x.head()
```

```
[12]:
```

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | B | LSTAT |
|---|---------|------|-------|------|-------|-------|------|--------|-----|-----|---------|--------|-----------|
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1 | 296 | 15.3 | 396.90 | 4.980000 |
| 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242 | 17.8 | 396.90 | 9.140000 |
| 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2 | 242 | 17.8 | 392.83 | 4.030000 |
| 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222 | 18.7 | 394.63 | 2.940000 |
| 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222 | 18.7 | 396.90 | 12.715432 |

```
[13]: y.head()
```

```
[13]: 0    24.0
      1    21.6
      2    34.7
      3    33.4
      4    36.2
      Name: PRICE, dtype: float64
```

```
[14]: from sklearn.model_selection import train_test_split
```

```
[15]: xtrain, xtest, ytrain, ytest=train_test_split(x,y,test_size=0.2, random_state=0)
```

```
[16]: xtrain.shape
```

```
[16]: (404, 13)
```

```
[17]: ytrain.shape
```

```
[17]: (404,)
```

```
[18]: xtest.shape
```

```
[18]: (102, 13)
```

```
[19]: ytest.shape
```

```
[19]: (102,)
```

```
[20]: from sklearn.linear_model import LinearRegression
```

```
[21]: # Fitting Multi Linear regression model to training model
      regressor=LinearRegression()
      regressor.fit(xtrain, ytrain)
```

```
[21]: LinearRegression
LinearRegression()
```

```
[22]: # predicting the test set results
ypred=regressor.predict(xtest)
```

```
[23]: from sklearn.metrics import mean_squared_error, mean_absolute_error
mse = mean_squared_error(ytest, ypred)
mae = mean_absolute_error(ytest,ypred)
print("Mean Square Error : ", mse)
print("Mean Absolute Error : ", mae)
```

```
Mean Square Error : 34.987389544238766
Mean Absolute Error : 3.9616211239591177
```

```
[24]: from sklearn.metrics import r2_score
```

```
[25]: r2=r2_score(ytest, ypred)
```

```
[26]: r2
```

```
[26]: 0.5703296053895559
```

```
[27]: # An R2 of 0.57 (57%) suggests a moderate fit of the model.
#The model is reasonably good at explaining the variability in the data,
#but there is still a significant portion (43%) that is not explained by the model
```

```
[ ]:
```